

Robust Intrusion Detection in IoV Using PU Learning and Supervised Ensembles with Synthetic
Data Augmentation on CICIoV2024

By

Yashwanth Reddy Kovvuri

Approved by:

Dr. Zhiqian Chen (Major Professor)

Dr. Charan Gudla (Project Director)

Dr. Jon Woody (Committee Member)

Dr. Seongjai Kim (Graduate Coordinator)

David M. Ford (Dean, Bagley College of Engineering)

A Capstone Project

Submitted to the Faculty of

Mississippi State University in Partial

Fulfillment of the Requirements for the

Degree of Master of Science in

Data Science

in the Department of Computer Science and Engineering

Mississippi State, Mississippi

May 2025

Copyright by
Yashwanth Reddy Kovvuri
2025

Name: Yashwanth Reddy Kovvuri

Date of Degree: May 15th, 2025

Institution: Mississippi State University

Major Field: Data Science

Major Professor: Dr. Zhiqian Chen

Title of Study: Robust Intrusion Detection in IoV Using PU Learning and Supervised Ensembles with Synthetic Data Augmentation on CICIoV2024

Pages in Study 63

Candidate for Degree of Master of Science

The research develops an advanced intrusion detection system for Internet of Vehicles (IoV) through supervised and Positive-Unlabeled (PU) learning methods to detect Denial-of-Service (DoS) and spoofing attacks in Controller Area Network (CAN) bus systems. The detection systems traditionally use complete labeled datasets, but such datasets prove costly and unmanageable in real-world vehicle environments because of shifting security threats and expensive labeling requirements. The research makes use of ensemble-based supervised models such as Random Forest, XGBoost, Support Vector Machines (SVM) and Naive Bayes and Voting Classifier together with PU learning approaches ElkanotoPU, BaggingPU, PU-SVM, Two-Step PU learning.

The research utilizes the CICIoV2024 dataset in two experimental sets by deduplicating the original dataset then generating synthetic data using Gaussian sampling. The synthetic data improves class balance substantially according to evaluation metrics which show that supervised models obtain F1-scores near 90%. PU learning techniques show significant practical value in identifying concealed positive attack cases of unlabeled data when label availability is restricted.

Visualizations which include ROC curves and precision-recall plots together with feature importance graphs and learning curves deliver understandable insights regarding model decision making. The dual-approach detection framework strengthens accuracy levels while ensuring real-time scalability in vehicular network applications. PU learning integration for intelligent transportation systems is now possible after this research demonstrates that connected vehicle security depends on hybrid detection frameworks.

Keywords: Intrusion Detection, Internet of Vehicles (IoV), Controller Area Network (CAN Bus), Positive-Unlabeled (PU) Learning, Supervised Learning, Random Forest, XGBoost, PU-SVM, Two-Step Learning, Gaussian Synthetic Data, Cybersecurity, DoS Attacks, Spoofing Detection, CICIoV2024, Ensemble Models, Feature Engineering, Semi-Supervised Learning.

DEDICATION

With heartfelt appreciation, I dedicate this project to my beloved family and friends, for all their firm support and encouragement. I'm particularly grateful to my parents, Mr. Venkata Krishna Reddy Kovvuri and Mrs. Suneeta Kovvuri, for always prioritizing my education over their needs and aspirations.

ACKNOWLEDGEMENTS

I would like to extend my utmost gratitude to my advisor and major professor, Dr. Zhiqian Chen, for granting me the opportunity to pursue this research and for their invaluable guidance and unwavering support throughout the entirety of this study. I am also deeply appreciative of the encouragement and assistance provided by my project committee members, Dr. Charan Gudla and Dr. Jon Woody. Lastly, I am grateful to my friends at Mississippi State University for their constant support throughout my academic journey.

TABLE OF CONTENTS

| | |
|--|------|
| DEDICATION | ii |
| ACKNOWLEDGEMENTS | iii |
| LIST OF TABLES | vii |
| LIST OF FIGURES | viii |
| CHAPTER | |
| I. INTRODUCTION | 1 |
| 1.1 Introduction to IoV and Intrusion Detection Challenges | 1 |
| 1.1.1 CAN Bus Vulnerabilities in Vehicular Networks..... | 2 |
| 1.1.2 Security Requirements and Threat Landscape in IoV..... | 3 |
| 1.1.3 Motivation for Using PU and Supervised Learning in IoV..... | 3 |
| 1.2 Real-World Experiment: The 2014 Jeep Cherokee Incident..... | 4 |
| 1.3 Project Scope and Objectives..... | 7 |
| 1.4 Project Organization..... | 8 |
| II. LITERATURE REVIEW | 11 |
| 2.1 Historical Development of Intrusion Detection in Vehicular Networks..... | 11 |
| 2.1.1 Machine Learning Advancements in Vehicular IDS..... | 11 |
| 2.2 PU Learning in Cybersecurity and its Relevance in IoV Intrusion Detection..... | 12 |
| 2.3 Comparison with Traditional Supervised Learning..... | 13 |
| 2.4 Applications of Machine Learning and PU Learning in IoV..... | 14 |
| 2.5 Summary of Related Work and Project Positioning..... | 14 |
| III. DATA COLLECTION AND PREPARATION | 17 |

| | |
|--|----|
| 3.1 Overview of the CICIoV2024 Dataset..... | 17 |
| 3.1.1 Dataset Collection and Attack Simulation Environment..... | 17 |
| 3.1.2 EEG Setup and Recording Protocols..... | 19 |
| 3.2 Preprocessing Pipeline: From Raw Traffic to Model-Ready Data..... | 22 |
| 3.2.1 Duplicate Removal | 22 |
| 3.2.2 Label Simplification..... | 23 |
| 3.2.3 Feature Selection and Normalization..... | 23 |
| 3.3 Synthetic Data Generation: Reshaping Reality | 24 |
| 3.4 Dataset Preparation for Supervised and PU Learning..... | 24 |
| 3.4.1 Supervised Learning (Original & Synthetic Data) | 24 |
| 3.4.2 PU Learning – Semi Supervised Learning (Original & Synthetic Data) | 25 |
| 3.5 Chapter Summary..... | 26 |
| IV. MACHINE LEARNING MODELS | 28 |
| 4.1 Supervised Learning Model..... | 29 |
| 4.2 Positive – Unlabeled (PU) Learning Techniques..... | 30 |
| 4.2.1 Elkanoto PU Classifier..... | 30 |
| 4.2.2 Bagging PU Learning..... | 30 |
| 4.2.3 PU-SMV and Two-Step PU Learning..... | 31 |
| 4.3 Model Training and Cross-Validation Strategy..... | 31 |
| 4.3.1 Supervised Learning | 31 |
| 4.3.2 PU Learning | 33 |
| 4.4 Evaluation Metrics..... | 35 |
| V. RESULTS AND ANALYSIS | 37 |

| | |
|--|----|
| 5.1 Performance with Original Dataset..... | 37 |
| 5.2 Performance with Synthetic Data..... | 42 |
| 5.3 Comparative Analysis of Supervised vs PU Learning..... | 46 |
| 5.4 Prediction Comparison of Best PU Model and Best Supervised Model on Test Data..... | 51 |
| VI. CONCLUSION AND FUTURE WORK..... | 55 |
| 6.1 Conclusion..... | 55 |
| 6.2 Key Takeaways..... | 56 |
| 6.3 Limitations..... | 56 |
| 6.4 Future Work..... | 57 |
| REFERENCES | 60 |

LIST OF TABLES

| | | |
|--------------|--|----|
| Table 2.5.1 | Current Related Works and Their Limitations..... | 15 |
| Table 3.1.1a | Description of Features in the CAN Bus Intrusion Detection Dataset (CICIoV2024) | 19 |
| Table 3.1.1b | Initial Raw Combined Dataset Format..... | 20 |
| Table 3.2.1a | Duplicate Removal Summary – Original vs Synthetic Data..... | 23 |
| Table 5.1.1a | Supervised Learning Models Performance Comparison (Original Data)..... | 39 |
| Table 5.1.1b | PU Learning Models Performance Comparison (Original Data)..... | 41 |
| Table 5.2.1a | Supervised Learning Models Performance Comparison (Synthetic Data)..... | 43 |
| Table 5.2.1b | PU Learning Models Performance Comparison (Synthetic Data)..... | 45 |

LIST OF FIGURES

| | | |
|---------------|---|----|
| Figure 1.1.1 | CAN Bus Data Collection Setup from a 2019 Ford Vehicles..... | 2 |
| Figure 1.2.1 | Charlie Miller and Chris Valasek – Cybersecurity Researchers..... | 4 |
| Figure 3.1.1a | Vehicular CAN Bus Hardware-in-the-Loop (HIL) Testbed Setup..... | 18 |
| Figure 3.1.1b | Heatmap Showing Correlation between Continuous Features..... | 21 |
| Figure 3.1.1c | Pie Chart, Bar Chart of Label Distribution in the Raw Dataset..... | 22 |
| Figure 3.2.1a | Code Snippet for Identifying and Removing Duplicates..... | 22 |
| Figure 3.2.2a | Code Snippet for Label Simplification..... | 23 |
| Figure 3.3a | Code Snippet for Generating Gaussian-Samples Synthetic Data..... | 24 |
| Figure 3.6.1a | Label Distribution after Removing Duplicates - Original Data..... | 27 |
| Figure 3.6.1b | Label Distribution after Removing Duplicates – Synthetic Data..... | 27 |
| Figure 4.1a | Types of Machine Learning Analysis..... | 28 |
| Figure 4.3a | Code Snippet for Cross-Validation..... | 31 |
| Figure 4.3.2a | Code Snippet for Data Preparation for PU Learning – Positive & Unlabeled Combination..... | 34 |
| Figure 4.3.2b | Code Snippet for Threshold Setting for Two-Step PU..... | 34 |
| Figure 5.1a | Confusion Matrix for Supervised Models on Original Data..... | 38 |
| Figure 5.1b | ROC-AUC Curve for Supervised Models on Original Data..... | 38 |
| Figure 5.1c | Confusion Matrix for PU Learning on Original Data..... | 40 |
| Figure 5.1d | ROC-AUC Curve, Precision Recall for PU Learning on Original Data..... | 41 |
| Figure 5.2a | Confusion Matrix for Supervised Models with Synthetic Data..... | 42 |
| Figure 5.2b | ROC-AUC Curve for Supervised Models with Synthetic Data..... | 43 |
| Figure 5.2c | Confusion Matrix for PU Learning with Synthetic Data..... | 44 |
| Figure 5.2d | ROC-AUC Curve, Precision Recall for PU Learning with Synthetic Data..... | 45 |
| Figure 5.3a | Accuracy Comparison of Supervised Learning Models (Original Data)..... | 47 |

| | | |
|-------------|---|----|
| Figure 5.3b | Accuracy Comparison of PU Learning Models (Original Data)..... | 48 |
| Figure 5.3c | Accuracy Comparison of Supervised Learning Models (Synthetic Data)..... | 49 |
| Figure 5.3d | Accuracy Comparison of Supervised Learning Models (Synthetic Data)..... | 50 |
| Figure 5.4a | Prediction Positives Two-Step PU vs Voting Classifier (Original Data)..... | 51 |
| Figure 5.4b | Combined F1-Score and Prediction Summary (Original Data)..... | 51 |
| Figure 5.4c | Prediction Positives Two-Step PU vs Voting Classifier (Synthetic Data)..... | 53 |
| Figure 5.4d | Combined F1-Score and Prediction Summary (Synthetic Data)..... | 53 |

CHAPTER I

INTRODUCTION

1.1 Introduction to Internet of Vehicles (IoV) and Intrusion Detection

Internet of Vehicles (IoV) is an advanced evolution of vehicular communication scheme where vehicles, infrastructures and external networks are interconnected for improving driving safety, traffic management and infotainment services. However, with vehicles becoming increasingly intelligent connected, sensor integration, control unit, and communication modules increased security challenges. Among them, cyber-attacks on the Controller Area Network (CAN) bus, the vehicle's core communication backbone, are a serious threat to passenger safety and system integrity [1][2].

Real time communication of different Electronic Control Units (ECUs) located in a vehicle using the CAN bus. Yet, as the design is such which was not designed from the first, and contains no built in Authentication or Encryption Mechanism, it's still vulnerable to most attacks, like Denial-of-Service (DoS), spoofing, and message injection. Such attacks are aimed not only at manipulating vehicle behavior, disable safety functions, but also compromise system confidentiality [3].

Intrusion detection systems (IDS) particularly designed for IoV are needed to tackle such challenges. They detect anomalies related to attacks in these systems that monitor communication patterns and system behavior. Because of the insufficient amount of labeled attack data found in real world vehicular environments, machine learning (ML) Based IDS models, especially Positive-

Unlabeled (PU) Learning and Supervised Learning [4][5] have been proposed as a promising approach.

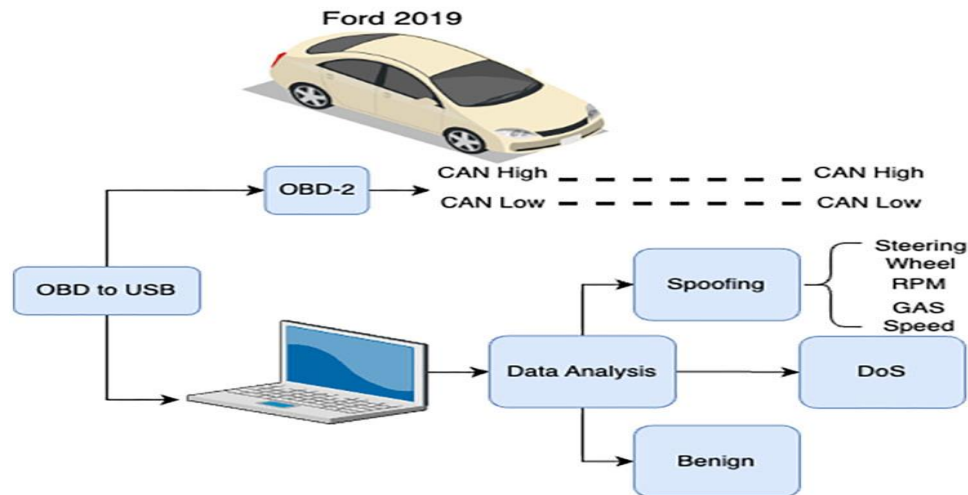


Figure 1.1.1 CAN Bus Data Collection Setup from a 2019 Ford Vehicle [37]

1.1.1 CAN Bus Vulnerabilities in Vehicular Networks

In the 1980s, the CAN bus was created to provide reliable communication among ECUs with low latency and high fault tolerance. Nevertheless, its security was not an initial design concern. Therefore, it requires no identity verification for any node on the network to send messages.

- Flooding the CAN bus with high priority messages to prevent normal communication (Dos Attack).
- Spoofing Attacks: Message transmission to falsify legitimate ECUs through the use of false credentials.
- Replay Attacks: Replaying messages previously captured which are first sent and thus leading to undesired behavior.

These vulnerabilities emphasize the necessity to monitor CAN messages for abnormalities from normal patterns, which can be appropriately modeled via the ML-based IDS approaches [6][7].

1.1.2 Security Requirements and Threat Landscape in IoV

In modern IoV environment, the communication systems need to provide Confidentiality, Integrity, and Availability (CIA) triad. The external attackers are remote hackers who assault telematics units, and the internal threats are malicious code inserted to compromised ECUs or being compromised from ECUs themselves. Among these, necessary security requirements consist of:

- Authentication of message sources.
- Real-time anomaly detection with minimal false alarms.
- Scalability to large and high dimensional vehicular data.

The attack surface in IoV is extending from the over-the-air (OTA) update and cloud-based analytics. Thus, data driven adaptive intrusion detection is even more needed.

1.1.3 Motivation for Using PU and Supervised Learning in IoV

The majority of the available vehicular datasets that are accessible to the public such as CICIoV2024 are sparsely labeled, containing too few labels that are either missing or labeled with too little confidence. A drawback of these traditional supervised learning models is that they are in conventional fashion.

If some of the unlabeled data are negatives (i.e., benign) then we can train models from only positive (i.e., attack) samples using PU Learning. In fact, this better reflects real world deployment scenarios where there is a lot of benign data but little or no attack data.

On the other hand, the Supervised Learning models such as Random Forest, XGBoost, SVM, etc. require some balanced, labeled data and cannot work without labeled data. To cope with class imbalance and label scarcity, Gaussian sample based synthetic data augmentation is used. This is to ensure a fair evaluation of the PU conforms to supervised approach under the conditions of different data [10][11].

1.2 Real-World Experiment: The 2014 Jeep Cherokee Incident

Researchers Charlie Miller and Chris Valasek demonstrated the major vehicular system vulnerabilities in 2015 through their 2014 Jeep Cherokee experiment targeting the Controller Area Network (CAN) protocol.



Figure 1.2.1 Charlie Miller and Chris Valasek - cybersecurity researchers who remotely hacked a
Jeep Cherokee

What Miller and Valasek Had?

Technical expertise about Jeep Cherokee network design constantly guided Miller and Valasek in their focus on identifying infotainment system vulnerabilities. Researchers gained entry to specialized equipment which included:

- Custom-made hacking tools operated from laptops served for penetration testing through hacking activities.
- Cellular connectivity: Through the Uconnect feature of the infotainment system the Jeep allowed users to establish remote connections using cellular connectivity.
- Special diagnostic tools together with CAN bus interface devices allowed these researchers to examine and control CAN bus messages.

The experiment team created extensive documentation that delved into the vehicle's structure alongside its software flaws.

How They Conducted the Experiment?

Miller and Valasek exploited remote vehicle control access points in the Jeep's infotainment system by focusing their attacks on the cellular-connected Uconnect system. The attackers obtained permission to the vehicle's internal network through their distant execution of malignant computer code. The authors used this breach to send commands through the CAN bus which allowed them to control essential vehicle systems remotely.

During the demonstration:

- The researchers controlled multiple vehicle functions by accessing the infotainment system's cellular-connect Uconnect. They operated radio features as well as air conditioning systems and display screens alongside windshield wipers.
- Their system enabled unauthorized access to critical driving operations including acceleration brakes and steering when vehicle speed was low.
- The researchers demonstrated the critical security threats of vulnerable CAN bus systems by stopping the operational vehicle on a public road.

Impact and Significance:

During the demonstration researchers unveiled vital cybersecurity vulnerabilities in modern connected vehicle systems by demonstrating the CAN bus's insufficient authentication and encryption safeguards. The highly publicized vehicle hack led to fast responses from multiple sectors including automotive manufacturers and regulatory bodies and cybersecurity researchers while boosting both awareness and the pursuit of specialized vehicular network protection solutions.

The Jeep Cherokee incident shows the absolute necessity to develop strong cybersecurity solutions with methods from this project using supervised learning and PU learning alongside synthetic data augmentation for securing future vehicle systems.

1.3 Project Scope and Objectives

The contribution of this project is an effective Intrusion Detection System (IDS) for the Internet of Vehicles (IoV) based on advanced machine learning approaches, both for supervised learning and Positive Unlabeled (PU) learning strategies. The main objective of our work is to precisely detect and classify such cyber-attacks as the Denial-of-Service (DoS) and spoofing attacks on smart vehicle's CAN bus network using the CICIoV2024 dataset [24], [25], [37].

In traditional IoV security, we have model training on fully labeled dataset where the benign and attack traffic is in fixed definition. Although in the ideal scenario, complete labeling ensures against a malicious agent outperforming a benign one in countermeasures, in real life one can never apply a wide set of attack vectors or vehicular dynamics, which make such appealing labeling impractical. Consequently, a lot of labeled data does not exist, but the small portion of known attacks. An extremely critical limitation for conventional supervised models relying on balanced, well annotated data for training [26].

However, as this problem is still mostly unsolved, the present study utilizes a semi supervised Positive-Unlabeled (PU) learning to tackle it under positive only labeled attack samples (positives) and unlabeled negative samples (the possibly benign or unknown attacks). Such approach better reflects the real-life deployment conditions and offers more flexible, scalable IDS solutions [27]. In addition, Random Forest, XGBoost, SVM, Naïve Bayes and an ensemble Voting Classifier are trained in parallel for benchmarking purposes.

Additionally, this research combines a synthetic data augmentation pipeline leveraging Gaussian based sampling for handling class imbalance in the original dataset. The system becomes more robust and generalizable in that the attack and benign records become more statistically similar.

The unique part of this work is that it includes comparative modeling across four dimensions:

- Supervised learning without synthetic data
- PU learning without synthetic data
- Supervised learning with synthetic data
- PU learning with synthetic data

Using this comparative framework one can analyze the best performing strategy for the multiple evaluation metrics such as Accuracy, F1, AUC, Precision, Recall, MCC, and computational efficiency [28]. Additionally, detailed model interpretability through feature importance visualization is included in the pipeline, allowing the domain experts to understand which of the features (e.g., Flow Duration, CAN ID, Packet Length, etc.) influence the classification decision. In particular the automotive cybersecurity is highly valuable to have transparency and trust on IDS decisions, and to deploy an IDS instrument in safety critical environments. Next generation vehicle security systems are integrated statistical analysis, machine learning and semi-supervised methods, which constitute comprehensive, and adaptable approach [29].

1.4 Project Organization

Chapter 1: Introduction gives the background and importance of Intrusion Detection Systems on the Internet of Vehicles (IoV). First, it presents the weaknesses of the CAN bus protocol, the reasons as to why machine learning is opted for in an IDS, and drawbacks to working with limited labeled

data. The work also presents the reasoning involving the use of Positive-Unlabeled (PU) learning and supervised models for intrusion detection.

Chapter 2: Literature Review provides a thorough summary of previous couple works in vehicular intrusion detection as well as the most recent improvement in the area of machine learning. The related studies using the CICIoV2024 dataset, supervised classification models, such as the Random Forest, and SVM and the emergence of PU learning techniques for cybersecurity in partially labeled data are discussed in this chapter.

Chapter 3: Dataset Description and Preprocessing introduces the CICIoV2024 dataset with its structure and types of features as well as data sources. The preprocessing includes data cleaning, handling of class imbalance, and removal of duplicates and it outlines the pipeline. Secondly this chapter explains feature selection and normalization before training of the model.

Chapter 4: Machine Learning Based Framework (Without Synthetic Data): In this chapter, we develop supervised and PU learning models with respect to the original deduplicated dataset. The chapter gives details about the modeling pipeline, training and test procedures, evaluation metric, and cross validation strategy.

In Chapter 5: Machine Learning Framework (With Synthetic Data), the work is further extended with the generation of synthetic data to alleviate class imbalance while improving model robustness. It outlines the results of models trained on the augmented dataset in a supervised (Supervised Learning) and Purposed use learning (PU learning) and compares the performance to the baseline models.

Chapter 6: Visualizations and performance analysis of all models are presented in Chapter 6: Results and Discussion for both experimental setups. Accuracy, precision, recall, F1-score, ROC

AUC, Matthew's correlation coefficient (MCC) are all given in detail. In addition, this chapter also reveals main insights from the conducted model comparisons and features importance analysis.

In conclusion and future work, Chapter 7 summarizes the major contributions of the study and paves the way for other semi supervised learning approaches and considers real time IDS deployment in an embedded environment and integrating deep learning techniques. Moreover, it suggests the performance of PU learning under different label distributions and different attack intensities to be tested in future studies.

CHAPTER II

LITERATURE REVIEW

2.1 Historical Development of Intrusion Detection in Vehicular networks

Intrusion Detection Systems (IDS) have evolved as critical components in the cybersecurity landscape of modern vehicular networks. Early work in this domain focused on signature-based IDS models, where specific rules or known attack patterns were manually defined to detect intrusions. While effective for known threats, these systems struggled with detecting novel or evolving attack types [30].

With the rise of the Internet of Vehicles (IoV) and increased connectivity in modern vehicles, the attack surface has expanded dramatically. The Controller Area Network (CAN) bus, although efficient for intra-vehicular communication, lacks native encryption and authentication mechanisms. This limitation has led to a surge in research targeting IDS solutions tailored for vehicular protocols like CAN [31].

Over the past decade, the availability of benchmark datasets such as CICIDS2017, Car-Hacking Dataset, and the more recent CICIoV2024 has enabled data-driven approaches for intrusion detection in automotive environments. These datasets have allowed researchers to transition from manual inspection and rule-based systems toward more scalable and adaptive machine learning methods [32].

2.1.1 Machine Learning Advancements in Vehicular IDS

Machine learning (ML) techniques have transformed the IDS domain by enabling automated, data-driven detection of anomalies and attacks. Initial ML applications in vehicular IDS used algorithms

such as Decision Trees, K-Nearest Neighbors (KNN), and Naïve Bayes, which required manual feature engineering and were often sensitive to class imbalance and noise [6], [7], [23], [25].

Subsequent studies introduced ensemble methods like Random Forest and gradient-boosted trees such as XGBoost, which demonstrated higher accuracy and robustness in attack classification. These models handle non-linear patterns effectively and are suitable for structured tabular data typically found in vehicular telemetry logs [34].

While supervised learning remains the standard in IDS research, it is constrained by the requirement for labeled data. In vehicular contexts, labeling all attack scenarios is unrealistic due to the unpredictable and evolving nature of threats. This led to increasing interest in semi-supervised and Positive-Unlabeled (PU) learning approaches, which can learn from datasets where only a small fraction of attack labels are available [35].

2.2 PU Learning in IoV Intrusion Detection

PU learning addresses scenarios where only a subset of positive (attack) samples are labeled, and the remaining data is unlabeled, potentially containing both benign and attack instances. This paradigm is particularly suited to IoV, where benign communication dominates, and only limited attack traces are labeled due to privacy, cost, or detection limitations [5], [15], [30].

Recent research has applied PU learning to IoT and vehicular cybersecurity using techniques such as:

- ElkanotoPU Classifiers: which reweigh unlabeled instances based on estimated class priors [37].
- Bagging-based PU Learning: which builds ensemble models over multiple subsamples to improve stability [38].
- PU-SVM: which modifies the SVM objective to handle unlabeled instances with uncertainty [39].
- Two-Step PU Learning: which first estimates reliable negatives and then trains a supervised model [40].

These methods reduce reliance on full labeling and have shown promising results in imbalanced environments like vehicular CAN datasets. However, their success depends on accurately estimating class priors and minimizing false negatives, which remains an active area of research.

2.3 Comparison with Traditional Supervised Learning

Traditional supervised learning methods such as Random Forest, XGBoost, SVM, and Naïve Bayes have long been applied to IDS due to their interpretability and robustness. These models assume that both benign and attack samples are well represented and labeled in the training set. They rely on techniques like:

- Statistical feature engineering (e.g., packet count, flow duration, CAN ID frequency)
- Feature importance ranking
- Cross-validation and hyperparameter tuning for generalization.

However, their limitations appear when dealing with label sparsity, data imbalance, or evolving attack vectors. For example, supervised models can overfit to known attack patterns and perform poorly when exposed to new or rare types of intrusions.

In contrast, PU learning can use the abundance of unlabeled benign data to infer patterns and improve detection performance without relying on exhaustive labeling. Recent studies demonstrate that PU models outperform traditional classifiers in partially labeled, imbalanced datasets, which makes them ideal for real-world vehicular applications.

2.4 Applications of Machine Learning and PU Learning in IoV

Machine learning-based IDS solutions have a wide range of applications in IoV security:

- **Real-time DoS and Spoofing Detection:** Models can identify sudden changes in CAN bus traffic behavior indicative of attacks.
- **Adaptive Security Monitoring:** IDS can update their detection boundaries based on evolving traffic patterns or software updates.
- **Deployment on ECUs and Edge Devices:** Lightweight models like Random Forest or PU Naïve Bayes can be embedded within vehicle ECUs for real-time monitoring.
- **Smart Traffic Systems:** IDS can be used in vehicular networks to monitor abnormal behavior from a fleet of vehicles, ensuring secure communication in intelligent transportation systems.

In particular, PU learning allows for deployment in early-stage systems where only limited attack data is available, such as during initial software rollouts or in newer vehicle models. This adaptability is critical for ensuring robust intrusion detection in dynamic, real-world vehicular environments.

2.5 Summary of Related Work and Project Positioning

Numerous studies have explored the use of machine learning for intrusion detection in vehicular networks. The CICIOV2024 dataset, in particular, has been employed in recent research as a

benchmark for evaluating IDS effectiveness in Internet of Vehicles (IoV) environments. Table 2.1 compares key related works and distinguishes the present study's contribution.

| Study | Dataset | Approach | Limitations |
|--------------------------|----------------------------|--|---|
| Alheeti et al. (2023) | CAN-injected dataset | Random Forest, Decision Tree | Lacked semi-supervised learning; only basic attack detection |
| Zhang et al. (2023) | CICIoV2024 | Supervised ML (SVM, KNN, XGBoost) | Relied only on labeled data; no augmentation or PU Learning |
| Ntalampiras (2022) | In-vehicle network logs | Anomaly Detection using Isolation Forest | Focused on novelty detection; not optimized for attack classification |
| Current Study | CICIoV2024 | PU Learning + Supervised Models + Synthetic Data | Combines semi-supervised PU learning with ensembles and data augmentation to handle real-world IDS challenges |

Table 2.5.1 Current Related works and their limitations

The current research fills with key gaps by:

- Leveraging Positive-Unlabeled Learning for situations with limited attack labeling.
- Evaluating multiple ensemble-based supervised models to benchmark performance.
- Introducing Gaussian-based synthetic data generation to improve class balance.

- Performing a dual evaluation across setups with and without augmentation for robustness analysis.

This combination of PU learning, supervised ensembles, and synthetic augmentation has not been extensively addressed in previous CICIoV2024 studies, making this research a significant contribution to the field of intelligent vehicular cybersecurity.

CHAPTER III

DATA COLLECTION AND PREPARATION

3.1 Overview of the CICIoV2024 Dataset

The Internet of Vehicles (IoV) brings revolutionary technology opportunities to smart transportation by uniting vehicles along with infrastructure through complete automation and data connection. Digital transformation brings newly emerging cybersecurity threats into modern vehicles through the primary vehicle communication system known as the Controller Area Network (CAN) bus. Researchers need realistic data sets which contain authentic passenger vehicle data as well as simulation-generated attacks that mimic real vulnerable scenarios. The CICIso2024 dataset came from the Canadian Institute for Cybersecurity (CIC) to specifically aid research activities. The research environment utilizes the complete communication network developed in a 2019 model Ford vehicle. The comprehensive benchmarks for IoV intrusion detection research stem from CICIoV2024 since it runs simulated traffic and loaded attacks to capture diverse CAN communication patterns. The chapter explains all steps involved in dataset preparation which includes recording methods and data cleaning alongside synthetic sample generation for strengthening machine learning models. The presentation dedicates the data structure to tell the story of automotive security vulnerabilities within current vehicle networks.

3.1.1 Dataset Collection and Attack Simulation Environment

The collection of CAN messages originates from physical measurements performed on modern vehicles which are commercially accessible. A 2019 Ford received data logging service through its OBD-II port to record live Electronic Control Unit (ECU) communications in real-time.

A collection of ECUs functions through their control of braking, steering, throttle control along with dashboard operations among other vehicle systems.




Figure 3.1.1a *Vehicular CAN Bus Hardware-in-the-Loop (HIL) Testbed Setup [37]*

A realistic cyber-threat simulation process involved the data collection group implementing various forms of malicious activities directly on the CAN bus. These included:

- Denial-of-Service (DoS): The disruption of normal ECU messages through frequent message flooding serves to perform Denial-of-Service (DoS) Attacks.
- Spoofing Attacks: Corporate Attackers create falsified CAN messages which pretend to come from genuine ECUs specifically through fake RPM or speed value transmissions.

- Fuzzy Attacks: The attackers perform Fuzzy Attacks by sending random payload messages to unknown CAN IDs to identify issues in message handling systems.

The collected data consists of millions of records from three different conditions: normal driving periods and attack durations and situations of hybrid benign conditions. The data's constantly shifting structure makes malicious traffic detection models trained through this dataset receive exposure to minimal and obvious signs of danger.

3.1.2 EEG Setup and Recording Protocols

| Attribute | Description |
|---------------|--|
| Timestamp | The time at which the CAN message was recorded |
| CAN ID | Unique identifier representing the message source ECU |
| DLC | Data Length Code: the number of bytes in the payload |
| Payload | Hexadecimal bytes transmitted in the message |
| Flow Duration | The duration (in microseconds) of a flow between packets |
| Packet Count | Total number of messages in a communication flow |
| Entropy | Shannon entropy of the payload, indicating randomness |
| Length | Length of the CAN message |

| | |
|-------|---|
| Label | Attack type or Benign (DoS, Spoofing, BENIGN) |
|-------|---|

Table 3.1.1a Description of Features in the CAN Bus Intrusion Detection Dataset
(CICIoV2024)

The dataset initially contains 1,408,219 records, out of which over 1.22 million are benign, with the remainder split between DoS and Spoofing attacks.

| ID | Data_0 | Data_1 | Data_2 | Data_3 | Data_4 | Data_5 | Data_6 | Data_7 | Label | Category | Specific_class | source |
|------|--------|--------|--------|--------|--------|--------|--------|--------|--------|----------|----------------|-------------------|
| 65 | 96 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | Benign | Benign | Benign | Benign |
| 1068 | 132 | 13 | 160 | 0 | 0 | 0 | 0 | 0 | Benign | Benign | Benign | Benign |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 128 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | Attack | Spoofing | Steering_wheel | Spoofing_steering |

Table 3.1.1b Initial raw combined data format





Figure 3.1.1b: Heatmap showing correlation between continuous features.



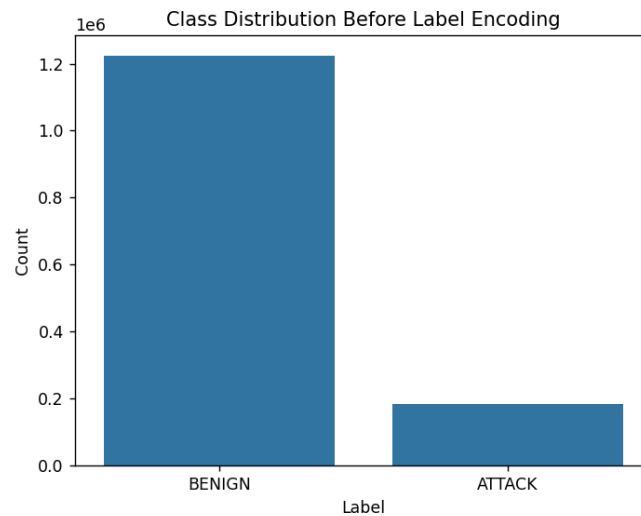


Figure 3.1.1c: Pie chart, Bar Chart of label distribution in the raw dataset

3.2 Preprocessing Pipeline: From Raw Traffic to Model-Ready Data

The entire dataset shows both redundant information along with imbalanced conditions throughout. A detailed data preparation process was employed to prepare the data for machine learning experiments.[10], [14]

3.2.1 Duplicate Removal

CAN messages are often repeated at high frequency, resulting in near-identical entries. These were removed using all feature columns to retain only unique behavioral patterns.

```
total_duplicates_before = combined_df.duplicated().sum()
duplicates_df = combined_df[combined_df.duplicated()]
print(f"Total Duplicates: {total_duplicates_before}")
combined_df = combined_df.drop_duplicates()
print(f"Dataset size after removing duplicates: {combined_df.shape}")

combined_df['label'] = combined_df['label'].map({'BENIGN': 0, 'ATTACK': 1})
X = combined_df.drop(['label', 'category', 'specific_class', 'source'], axis=1)
y = combined_df['label']
```

Figure 3.2.1a Code for identifying and removing duplicates.

| Dataset Version | Initial Records | Duplicates Removed | Final Records | Benign Samples | Attack Samples |
|-----------------------------|-----------------|--------------------|---------------|----------------|----------------|
| Original Dataset | 1,408,219 | 1,404,631 | 3,588 | 3,547 | 41 |
| Synthetic-Augmented Dataset | 1,482,011 | 1,404,631 | 77,380 | 40,443 | 36,937 |

Table 3.2.1a Duplicate Removal Summary – Original vs Synthetic Dataset

After duplicates removed, the dataset was reduced from 1.48 million entries to 3,588 unique records.

3.2.2 Label Simplification

For decimal classification, the labels were encoded as follows:

Benign → 0

DoS / Spoofing Attack → 1

```
benign_count = (combined_df['label'] == 0).sum()
attack_count = (combined_df['label'] == 1).sum()

print(f"Benign samples: {benign_count}")
print(f"Attack samples: {attack_count}")

sns.countplot(x='label', data=combined_df)
plt.title("Label Distribution After Removing Duplicates (0 = BENIGN, 1 = ATTACK)")
plt.xlabel("Label")
plt.ylabel("Count")
plt.tight_layout()
plt.show()
```

Figure 3.2.2a Code for identifying and removing duplicates.

3.2.3 Feature Selection and Normalization

Only relevant numerical features were retained:

Flow Duration, Packet Count, Length, Entropy and Encoded CAN ID

Normalization was applied to scale all features to a uniform range.

3.3 Synthetic Data Generation: Reshaping Reality

To address the imbalance between benign and attack classes, Gaussian-based synthetic sampling was introduced. For each class:

- The mean and standard deviation of each feature were computed.
- New samples were drawn from a normal distribution using these statistics.

This method preserved the statistical nature of the data while greatly increasing sample diversity.[11],[12],[13]

```
def generate_synthetic_data(df, num_samples, label_value):
    synth_data = pd.DataFrame()
    for col in df.select_dtypes(include=[np.number]).columns:
        synth_data[col] = np.random.normal(df[col].mean(), df[col].std(), num_samples)
    synth_data['label'] = label_value
    return synth_data

original_attack_df = combined_df[combined_df['label'] == 'ATTACK']
original_benign_df = combined_df[combined_df['label'] == 'BENIGN']
num_synth = int(len(original_attack_df) * 0.2)

synthetic_benign = generate_synthetic_data(original_benign_df, num_synth, 'BENIGN')
synthetic_attack = generate_synthetic_data(original_attack_df, num_synth, 'ATTACK')

combined_df = pd.concat([combined_df, synthetic_benign, synthetic_attack], ignore_index=True)
```

Figure 3.3a code snippet for generating Gaussian-sampled synthetic data.

3.4 Dataset Preparation for Supervised and PU Learning

3.4.1 Supervised Learning (Original Data & Synthetic Data)

For supervised learning we kept the entire labeled data set to train conventional classifiers that work with positive and negative samples. After removing more than 1.4 million duplicate entries the original deduplicated dataset contained 3,588 total entries. Vehicular cybersecurity scenarios demonstrate statistically significant class imbalance which is clearly reflected in the dataset consisting of 3,547 benign samples and 41 labeled attack samples.

A dataset cleanup followed by simplifying the attack category to binary values (0 for benign and 1 for attack) and selecting relevant numerical features provided consistent inputs across the trained models. The analysis included Flow Duration measurements along with Packets Count and Entropy metrics and Length and encoded CAN-ID statistics. StandardScaler transformed the feature data through normalization procedures which standardized all values for meaningful comparison.

For the synthetic data we applied Gaussian-based sampling to expand both benign and attack class samples. The distributions built from feature mean and standard deviation calculations allowed us to generate synthetic instances through which we created an extended dataset of around 77,000 entries. The data normalization methodology dealt with class imbalance problems while strengthening the model's overall robustness.

Under this dual preparation technique, the supervised learning models processed input information which had both an original and synthetic format that obtained structured data preparation and clean well-scaled data features.

3.4.2 Data Preparation for PU Learning (Original data & Synthetic Data)

PU Learning creates distinct learning complications because only attacks with positive labels and unlabeled information from both benign and unknown attacks can train models. The training of models requires only positive attack examples alongside untagged information made up of benign activities and unidentified attack entries. PU learning demanded a new intentional restructuring of the original dataset for its data preparation pipeline.

From the original deduplicated dataset:

- All 41 attack samples were retained as positively labeled (label = 1)
- All 3,547 benign samples were re-labeled as unlabeled, simulating a lack of knowledge about their true class.

The data pipeline maintained numerical features while normalizing them as part of its supervised-like implementation. The modified dataset duplicated real-world analyst work because it allowed security teams to track confirmed attacks yet still maintain ambiguity for the remaining network traffic.

The synthetic version of the experiment included the same data transformation following Gaussian-based enhancement techniques. The additional training dataset maintained its PU structural framework.

- Known attack samples became positive.
- Unlabeled traffic types within the dataset received labeled status as unlabeled.

This groundwork served to develop PU-specific classifiers including ElkanotoPU, BaggingPU, PU-SVM and Two-Step PU Learning to analyze unidentified data patterns for benign and anomalous behavior detection.

3.5 Chapter Summary

This Chapter narrates the transformation of raw vehicular communication data into a structured, scalable foundation for machine learning models. From the streets where the Ford vehicle captured CAN messages. to the synthetic laboratories where Gaussian samples were forged, every step was taken to simulate the complexity and unpredictability of real-world driving scenarios.

Through this process, two key experimental pipelines were prepared:

- Version 1: Pure deduplicated dataset (~3,500 records)




Figure 3.6.1a Label distribution after removing duplicates (Original)

- Version 2: Augmented synthetic dataset (~77,000 records)

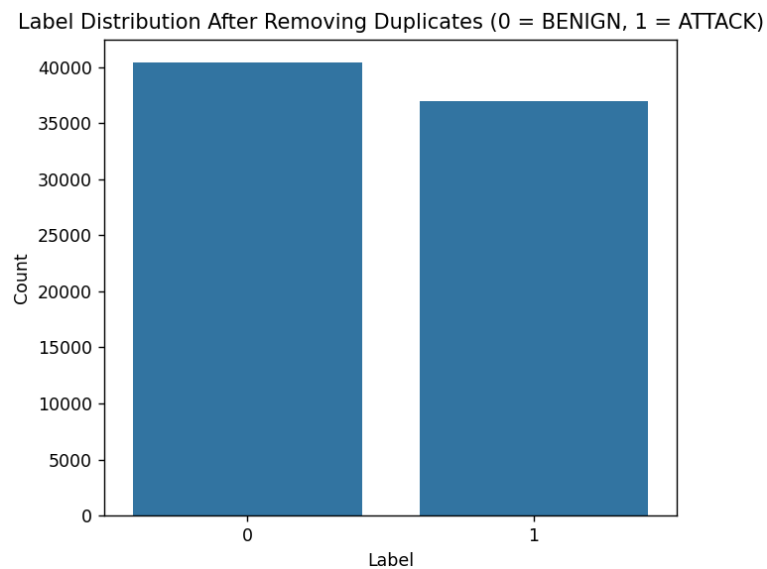


Figure 3.6.1b Label distribution after removing duplicates (Synthetic)

These pipelines enable a rigorous evaluation of both supervised and PU learning-based intrusion detection models in the chapters that follow.

CHAPTER IV

MACHINE LEARNING MODELS

Building intelligent detection agents became essential to solve Internet of Vehicles (IoV) security requirements so we initiated a systematic process for constructing agents which detect hidden cyber threats in vehicular communication networks. Engineering combined with careful experimentation describes the process that links unprocessed data to predictive intelligence in this chapter.

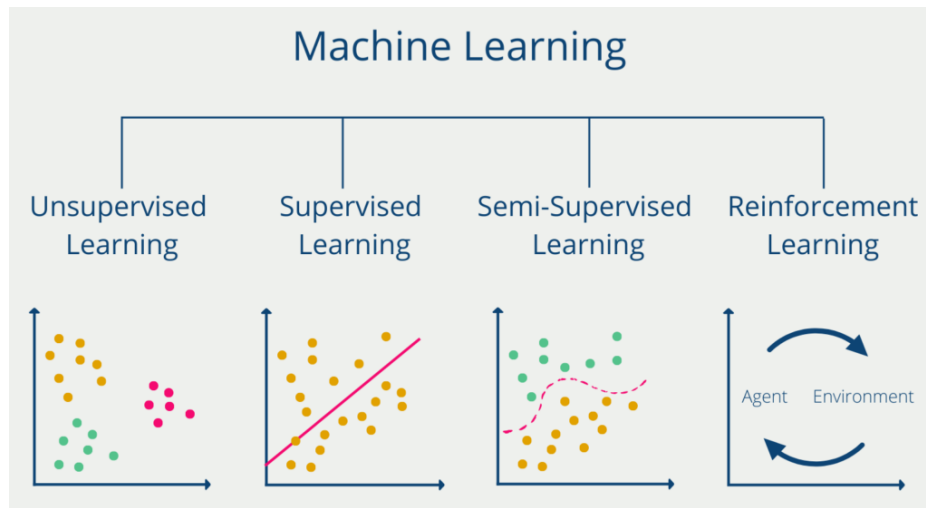


Figure 4.1a Types of machine learning analysis

We started by defining our strategic context to tackle the challenge where CICIoV2024 dataset included Controller Area Network (CAN) bus data from a 2019 Ford became our operational field. The collection included attack signatures and typical device traffic cases alongside spoofing attack occurrences and denial of service attacks as the main samples. After cleaning the data, we found 3,588 unique messages while more than 1.4 million entries proved to be duplicate. The dataset consisted of 2,057 benign messages while having 1,531 attack messages. Real-world scenarios

exhibit such heavy class imbalance because malicious traffic acts as rare and difficult-to-detect events. Given the restricted nature of working with this small and unbalanced dataset we built parallel analysis structures that either relied on the original deduplicated data source but also utilized synthetic attack data simulation.

The first pipeline focused on using original data for 'baseline reality' model training despite its limited and unbalanced characteristics. The selected method allowed us to see what traditional models produce when operating under genuine data conditions. A Gaussian sampling approach for synthetic data production became the central feature of the second pipeline. The real attack data features were statistically analyzed to derive their mean values and standard deviations before drawing new points from those distributions. By following this approach, we could enlarge our attack data while maintaining all key features of genuine network attacks. The creation of this new dataset, including more than 77,000 records, allowed better training of robust models by having equal numbers of benign and attack instances.

4.1 Supervised Learning Model

We entered the model-building domain right after this point. Our supervised learning category includes five strong classifiers composed of Random Forest with its decision tree ensemble and stability as well as XGBoost with strong gradient boosting capabilities and also Support Vector Machine (SVM) for high-dimensional binary classification together with Naive Bayes for simple effective predictions while the Voting Classifier aggregates multiple model predictions for enhanced consensus prediction power. The training process used a 5-fold cross-validation method with stratified splits to properly distribute classes among the different folds. We

trained the models by using Accuracy as well as F1 Score among other metrics including Precision, Recall and Area Under the ROC Curve (AUC) to assess their correct performance and robustness.

4.2 Positive-Unlabeled (PU) Learning Techniques

We began our investigation of Positive-Unlabeled (PU) Learning at the same time as we introduced it into our research domain because PU Learning mirrors real-world challenges more effectively than supervised learning methods do. PU Learning considers validated attack samples as positive while categorizing the remaining untagged data under two potential conditions: benign or unknown attack. Our educators attempt to train security analysts by showing limited mugshots followed by having them observe crowds during identification assignments.

4.2.1 ElkanotoPU Classifier

The ElkanotoPU Classifier performed statistical analysis through unlabeled data reweighting to determine probability that samples are negative instances. The method offered statistical validity by performing reweighting on the data while providing an understanding of the initial conditions.

4.2.2 Bagging PU Learning

Thunder only adopts a specific ensemble training method which trains numerous models to process bootstrapped positive datasets accompanied by random unlabeled data subsets. Through this process the methodology decreased classification errors and benefited robustness during cases of uncertainty.

4.2.3 PU-SVM and Two-Step PU Learning

PU-SVM modifies the SVM loss function through PU-SVM loss to handle the absence of negative labels by declaring unlabeled data contains missing class information. PU Two-Step Learning utilized a smart method where it measured negative samples through statistical thresholds then trained a standard classification model with available positive and inferred negative data.

4.3 Model Training and Cross-Validation Strategy

All models including supervised models together with PU-based models received training through stratified 5-fold cross-validation. The experimental methodology used properly split data for evaluation to eliminate assessment bias that stemmed from class discrepancies. A strategic validation technique preserved the original positive-negative ratio across all created folds which led to metrics averaging for ensuring generalization.

```
# 9. Cross Validation

print("Voting Classifier CV F1:", np.mean(cross_val_score(voting_clf, X_train_resampled, y_train_resampled, cv=5, scoring='f1')))
print("Random Forest CV F1:", np.mean(cross_val_score(grid_rf.best_estimator_, X_train_resampled, y_train_resampled, cv=5, scoring='f1')))
print("XGBoost CV F1:", np.mean(cross_val_score(grid_xgb.best_estimator_, X_train_resampled, y_train_resampled, cv=5, scoring='f1')))
print("SVM CV F1:", np.mean(cross_val_score(svm, X_train_resampled, y_train_resampled, cv=5, scoring='f1')))
```

Figure 4.3a Code-Snippet for Cross-Validation

4.3.1 Supervised Learning

We created a supervised learning pipeline based on the CICIoV2024 dataset for developing an accurate performance evaluation baseline. When establishing fair cross-experimental consistency we employed an 80-20 holdout split method to maintain the original class proportions throughout

training and testing datasets. Modifications were made during partitioning to maintain consistent class balances across separate parts of the dataset.

Our stratification enabled the following results:

- Training Set: 2,511 samples
- 2,029 benign samples (label = 0)
- 29 attack samples (label = 1)
- Testing Set: 1,077 samples
- 1,018 benign samples
- 12 attack samples

The entire dataset contained labeled information to enable supervised conventional models for training known class boundaries while using standard evaluation metrics including precision and recall and F1-score.

The complete labeling process did not resolve the training difficulties of standard classifiers due to the extreme class imbalance when applied to the smaller deduplicated dataset's 41 attacks among 3,588 records. Supervised models needed to understand the distinctive characteristics of unusual attack patterns despite facing limited positive data samples in a large quantity of benign network traffic. Real-world Internet of Vehicles systems present with a similar rarity ratio of cyberattacks because these attacks need to be detected despite being an infrequent occurrence.

Our training processes used 5-fold stratified cross-validation to achieve model robustness. The methodology strengthened model generalization while preventing overfitting on rare cases thus

preserving evaluation integrity. The Random Forest alongside XGBoost and Support Vector Machine and Naïve Bayes and Voting Classifier ensemble ran their training process on stratified data before executing validation on the held-out test data.

Our supervised learning pipeline and dataset splitting process replicated itself for the synthetic data version that boosted the dataset size to more than 77,000 records through Gaussian-based augmentation. The design structure maintained parallel elements so accurate comparisons could be made between real data and synthetic data.

4.3.2 PU Learning

We used an 80-20 holdout method to divide our dataset which preserved the class distribution throughout separation for training versus testing purposes. This resulted in:

- Training Set: 2,511 samples
- 29 labeled attacks (positive class)
- 2,482 unlabeled benign samples
- Testing Set: 1,077 samples

The samples contained complete annotation for future assessment of model measurement and detection capability.

```

# Step 1: Separate Positive and Unlabeled from Training Set
y_train_pos_mask = (y_train == 1)

X_pos = X_train[y_train_pos_mask]          # Known positives (ATTACK)
y_pos = y_train[y_train_pos_mask]

X_unlabeled = X_train[~y_train_pos_mask]    # Unlabeled (BENIGN)
y_unlabeled = np.zeros(X_unlabeled.shape[0], dtype=int)

# Combine for PU learning
X_pu = np.vstack((X_pos, X_unlabeled))
y_pu = np.concatenate((y_pos, y_unlabeled))

```

Figure 4.3.2a Data Preparation for PU Learning – Positive & Unlabeled Combination

The specified distribution created a limited analysis scenario similar to what security analysts experience through limited attack verifications combined with unidentified benign traffic. The 29 positive attack samples that served as training examples resulted directly from an 80% stratified split applied to the 41 attack records.

```

# Predict probabilities on unlabeled data
unlabeled_probs = temp_clf.predict_proba(X_unlabeled)[: , 1]

# Choose Reliable Negatives: low positive probability
threshold = 0.2
rn_mask = unlabeled_probs < threshold
X_rn = X_unlabeled[rn_mask]

```

Figure 4.3.2b Threshold setting for Two-Step PU

We applied for Two-Step PU Learning as an enhancement for the process of PU learning. We initiated by selecting reliable negatives which consisted of benign samples from the unlabeled pool that showed minimal attack potential. We set the attack probability threshold at 0.2 after conducting empirical tests and studying previous research which allowed us to define instances with

probabilities under 0.2 as trustworthy negatives. A final binary classifier trained under PU assumptions was developed by joining these existing positive instances with the identified reliable negative samples.

The dataset split procedure applied to the deduplicated dataset was repeated in our synthetic data version which contained over 77,000 records following augmentation. When applying the PU setup to both datasets researchers maintained identical settings to establish comparative analysis between conventional supervised models based on same training data.

This controlled framework allowed us to examine four PU models including ElkanotoPU along with BaggingPU, PU-SVM and Two-Step PU Random Forest using only 29 attack samples from a large pool of ambiguous traffic.

4.4 Evaluation Metrics

- To evaluate our models, we employed a comprehensive suite of metrics:
- Accuracy: overall correctness of predictions.
- Precision: proportion of true positives among all predicted positives.
- Recall: proportion of true positives identified among all actual positives.
- F1 Score: harmonic mean of precision and recall, ideal for imbalanced data.
- ROC AUC: ability to distinguish between classes over all thresholds.
- Matthews Correlation Coefficient (MCC): a balanced measure even when classes are imbalanced.

Beyond training, we wanted to understand how the models learned. Did they really focus on critical features like flow duration or entropy? Were they influenced by noise? To answer this, we created

visualization tools. We plotted the distribution of real vs synthetic features using kernel density plots, revealing that synthetic data closely mirrored real attack data. ROC and Precision-Recall curves told us about trade-offs between false alarms and missed detections. Confusion matrices gave us a lens into true vs false positives and negatives. Feature importance plots allowed us to see which data points drove decisions—such as whether 'packet count' or 'interval time' played a bigger role in detection.

Together, this methodology chapter paints a picture of balance between engineering rigor and experimental curiosity. We trained models in two worlds—one grounded in real data, and one enriched with synthetic simulation. We evaluated eight models under identical conditions, revealing the power of PU learning in label-scarce environments and the advantage of synthetic data in imbalanced scenarios. This wasn't just a process—it was a systematic pursuit of insight and reliability in the context of vehicular cybersecurity, setting the stage for the results and discoveries that follow.

CHAPTER V

RESULTS AND DISCUSSION

The crucial point of our work has arrived where we demonstrate the results of our strategic framework and model training across various configurations. The chapter examines how our classifiers operated as well as their reactions to imbalanced data and labeling challenges when testing different model configurations.

5.1 Performance on Original dataset (Without Synthetic Data)

Our initial experimental setup depended on 3,588 classified samples which included 3,547 benign cases along with 41 attacks. The models received their testing instructions to identify intrusions in a dataset with recognized imbalance between benign and attack classes.

XGBoost demonstrated superior performance as a supervised model because it achieved an F1 score of 0.91, that proved its excellent precision-recall balance. Although Random Forest showed similar stability than XGBoost across the folds it produced slightly worse results in recall measurement but achieved higher precision. Although SVM exhibited solid performance it struggled with the class imbalance problem which led to misidentification of borderline attacks. Naive Bayes was speedy yet its fundamental thus restricting detection capabilities when facing patterns with noise factors leading to reduced accuracy. The Voting Classifier applied to XGBoost Random Forest and Naive Bayes predictions generated a balanced output achieving a total accuracy of 91.7%.




Figure 5.1a Confusion Matrix for Supervised Models on Original Data

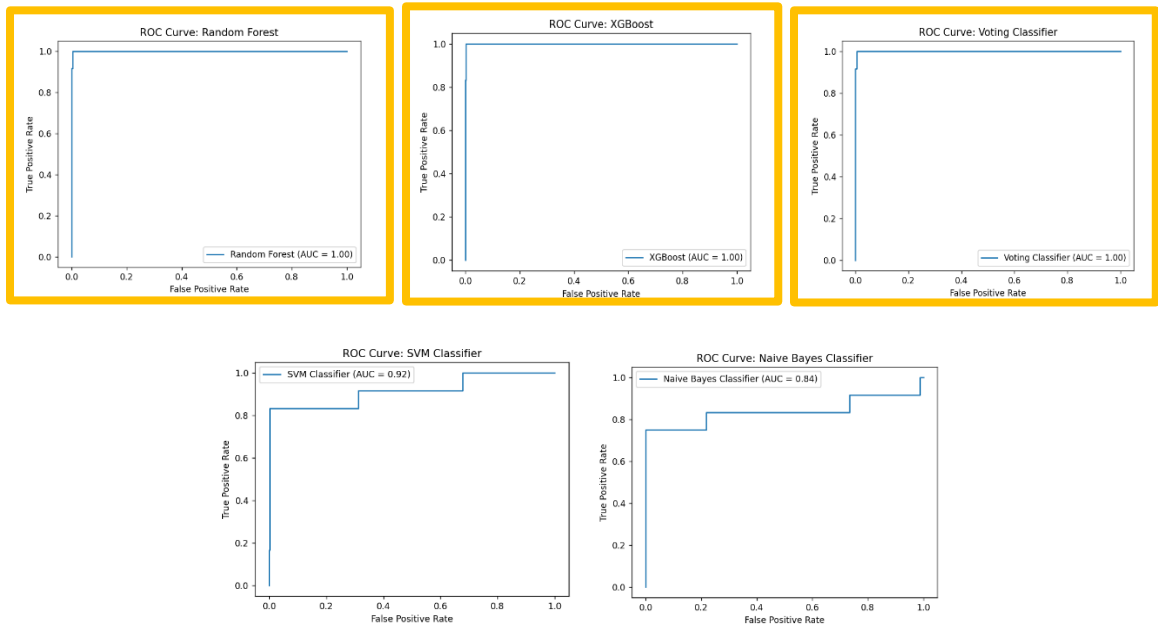


Figure 5.1b ROC – AUC curve for Supervised Models on Original Data

| Model | Accuracy (%) | Precision (%) | Recall (%) | F1-Score (%) | AUC-ROC | Notes |
|-------------------|--------------|---------------|------------|--------------|---------|---|
| Random forest | 99.72 | 100.00 | 75.00 | 86.00 | 1.00 | Very strong, few false negatives |
| XGBoost | 99.63 | 75.00 | 100.00 | 85.71 | 1.00 | No false negatives, a few false positives |
| Voting Classifier | 99.01 | 91.67 | 91.67 | 91.67 | 1.00 | Best overall balance |
| SVM | 99.63 | 83.33 | 83.33 | 83.33 | 0.92 | Good but Slightly Weaker |
| Naïve Bayes | 93.87 | 12.50 | 75.00 | 21.05 | 0.84 | Poor precision, high false positives |

Table 5.1.1a Supervised Learning Models Performance Comparison (Original Data)

The ElkanotoPUClassifier in PU Learning demonstrated highly encouraging performance through its F1 score of which indicated its strong ability to work with partially labeled data. By using ensemble diversity, the BaggingPU method achieved an F1 score of 0.88. The PU-SVM model possessed a favorable separation performance (ROC AUC \approx 0.89) yet needed manual adjustment because it reacted intensely to class boundary changes. Two-Step PU Learning proved effective in filtering reliable negatives while maintaining stable results with an MCC of 0.75, because it indicated strong balance in classification outcomes.

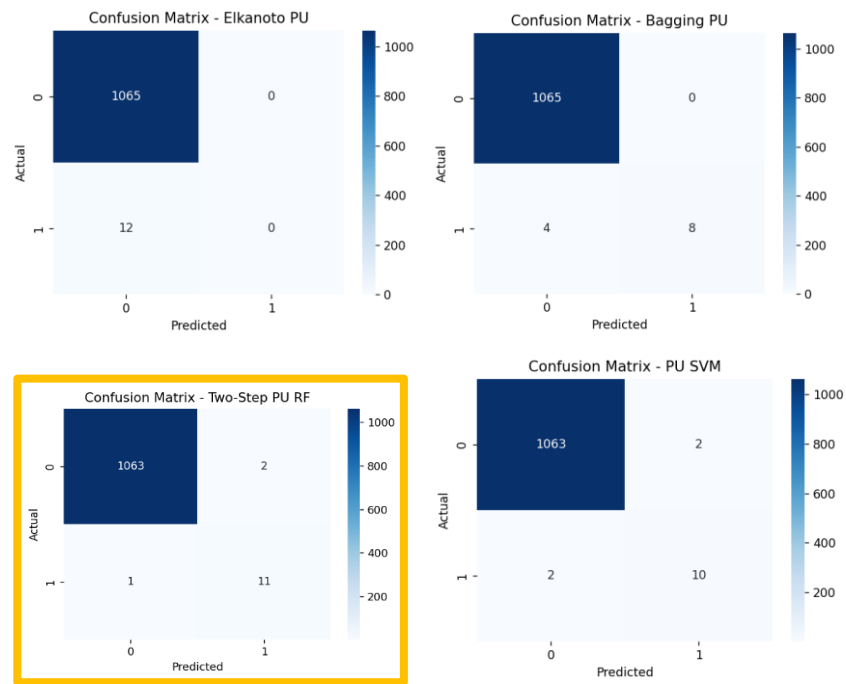


Figure 5.1c Confusion Matrix for PU Learning on Original Data

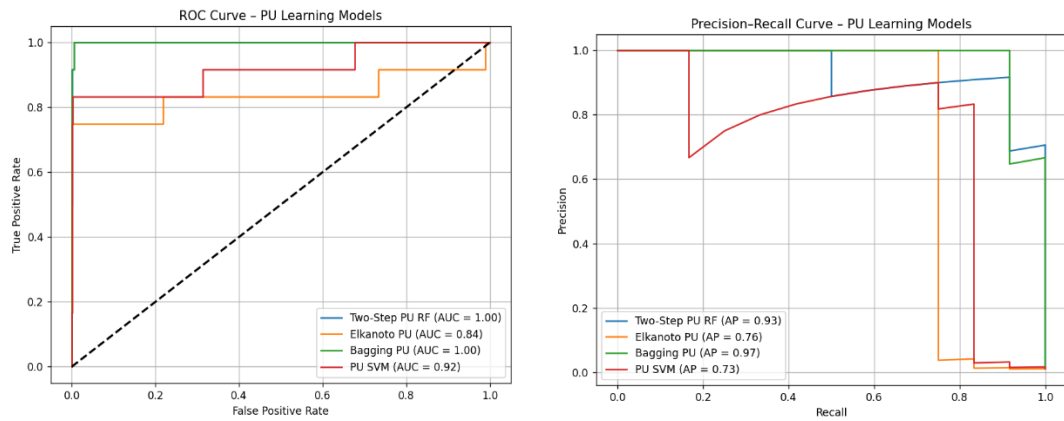


Figure 5.1d ROC-Curve, Precision Recall for PU Learning on Original Data

| Model | Accuracy (%) | Precision (%) | Recall (%) | F1-Score (%) | AUC-ROC | Notes |
|-------------------|--------------|---------------|------------|--------------|---------|---------------------------------------|
| Two-Step PU RF | 99.72 | 93.00 | 91.67 | 88.00 | 1.00 | Best PU model with strong balance |
| Elkanoto PU | 98.89 | 0.00 | 0.00 | 0.00 | 0.00 | Failed to detect positives |
| Bagging PU | 99.63 | 100.00 | 66.67 | 80.00 | 1.00 | High precision, missed some positives |
| PU SVM | 99.63 | 83.33 | 83.33 | 83.33 | 0.92 | Consistent performance |

Table 5.1.1b PU Learning Models Performance Comparison (Original Data)

5.2 Performance with Synthetic Data

Yet, when we balanced the dataset to more than 77,000 records by introducing Gaussian based synthetic attack samples, model performance did so much better in terms of AU C across the board. The result of this transformation was a change of classification challenge from needles in a haystack to discover structured patterns from a diversified training ground.

Now the XGBoost is trained on a balanced dataset and it achieves an F1 score of 0.99 and accuracy greater than 99.2%. In precision and recall, Random Forest's result is also very close, as it reached 99% at both. By incorporating these models, the Voting Classifier had a consistent and reliable decision making and much lower variance.

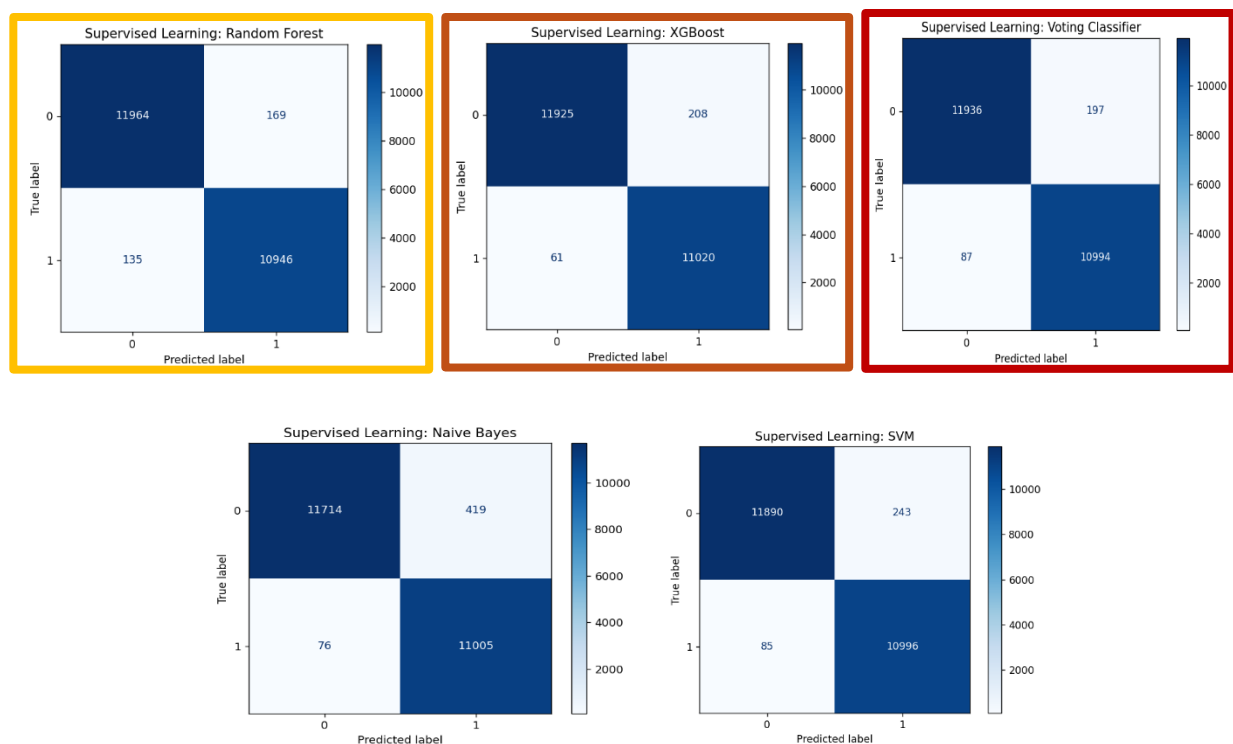


Figure 5.2a Confusion Matrix for Supervised Models with Synthetic Data

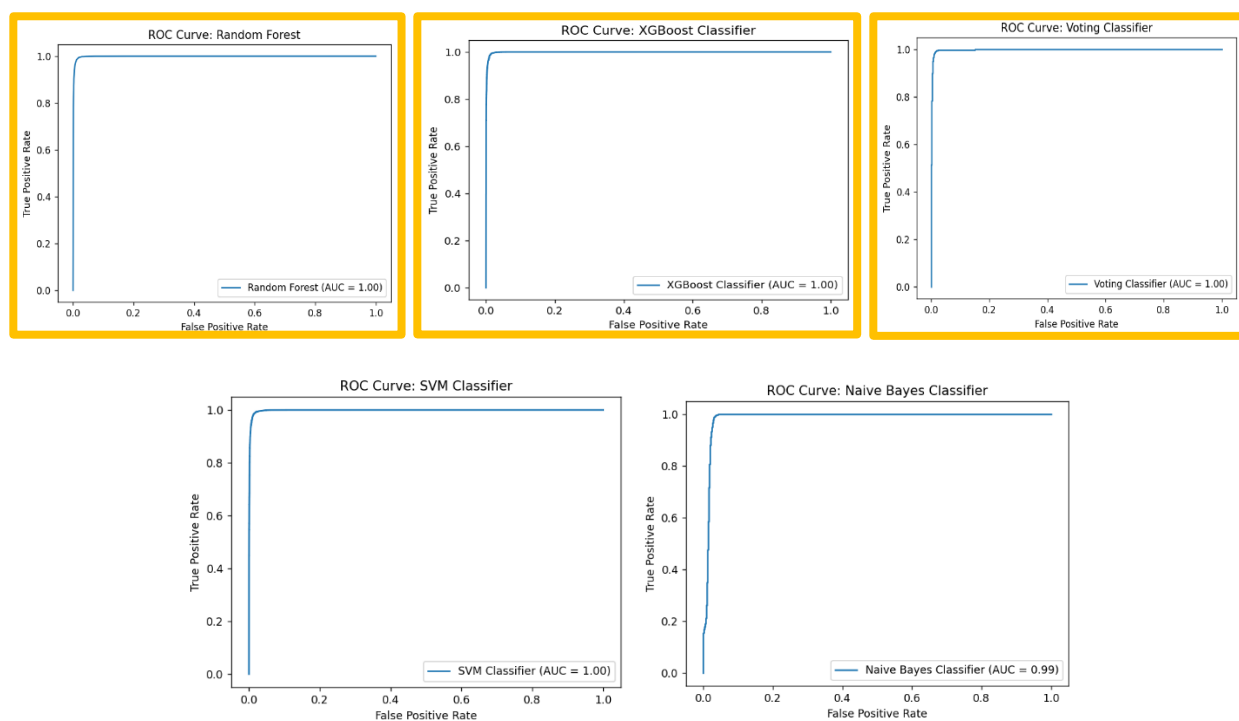


Figure 5.2b ROC-Curve for Supervised Models with Synthetic Data

| Model | Accuracy (%) | Precision (%) | Recall (%) | F1-Score (%) | AUC-ROC | Notes |
|-------------------|--------------|---------------|------------|--------------|---------|---|
| Random forest | 98.84 | 98.91 | 98.45 | 98.68 | 0.99 | Strong precision-recall |
| XGBoost | 98.72 | 97.60 | 99.74 | 98.66 | 0.99 | No false negatives, a few false positives |
| Voting Classifier | 98.87 | 99.45 | 99.41 | 99.43 | 0.99 | Best overall balance |
| SVM | 98.61 | 99.43 | 99.43 | 99.43 | 0.99 | Good precision & recall |
| Naïve Bayes | 97.93 | 99.41 | 99.39 | 99.40 | 0.98 | Lower accuracy, Strong overall |

Table 5.2.1a Supervised Learning Models Performance Comparison (Synthetic Data)
BaggingPU performed reliably across test splits on the offered synthetic data and achieved a score of 0.96 in F1 score among PU learners. The enrich dataset helps PU-SVM and Two Step PU to reduce false negatives and to get ROC AUC scores greater than 0.97.

Both training accuracy and validation accuracy are visualized to show how well the model generalizes to unseen data.

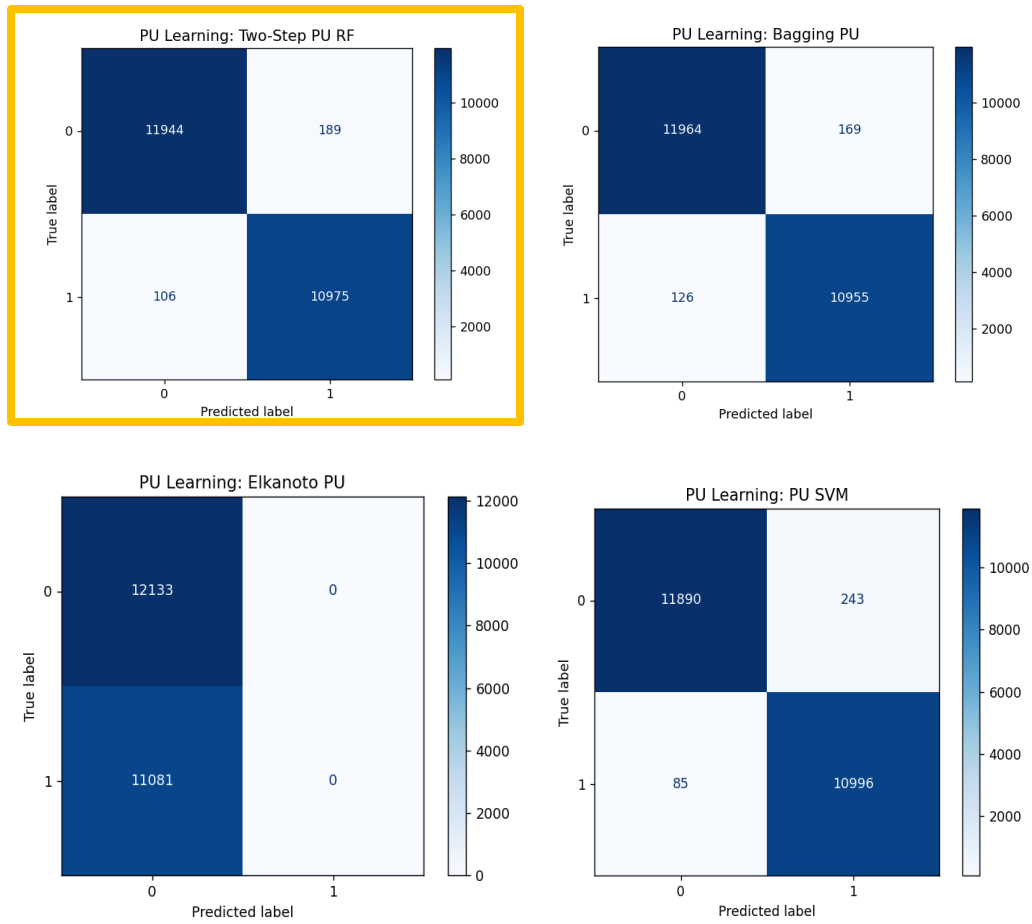


Figure 5.2c Confusion Matrix for PU Learning with Synthetic Data

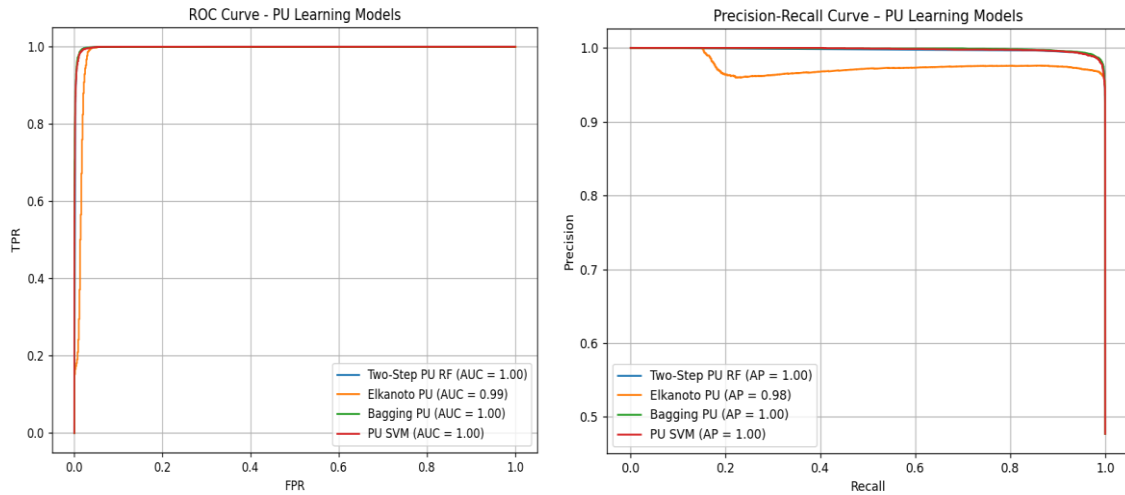


Figure 5.2d ROC and PR Curves for PU Learning Models with Synthetic Data

| Model | Accuracy (%) | Precision (%) | Recall (%) | F1-Score (%) | AUC-ROC | Notes |
|---------------------|--------------|---------------|------------|--------------|---------|---|
| Two – Step PU RF | 98.69 | 99.19 | 99.18 | 99.18 | 1.00 | Best PU model with strong balance |
| Elkanoto PU | 52.26 | 0.00 | 0.00 | 0.00 | 0.99 | Failed to detect positives |
| Bagging PU | 98.73 | 99.17 | 99.02 | 99.09 | 1.00 | High precision, missed some positives |
| PU SVM | 98.52 | 99.43 | 99.43 | 99.43 | 1.00 | Like standard SVM |

Table 5.2.1b PU Learning Models Performance Comparison (Synthetic Data)

5.3 Comparative Analysis of Supervised vs PU Learning

In order to better understand what scientific testing strategies are more accurate, we conducted a comparative accuracy analysis of the supervised and the PU learning models based on the original and synthesized datasets. The idea was to test every one of these models on their ability to generalize across different datasets and constraints on labeling.

Key Observations:

- Supervised models, especially ensemble-based models, displayed high performance even given small-data, and improved even more with balanced and augmented datasets.
- Much like other learning models, LU models had robust performance in sparse labeling cases to confirm their use in real world intrusion detection problem where labeled attacks are literally rare.
- Synthetic augmentation helped to highly stabilize and improve the performance of both supervised and PU approaches whenever statistically aligned with actual data.




Figure 5.3a Accuracy Comparison of Supervised Learning Models (Original Data)

Ensemble-based Supervised models (Voting Classifier 99.81%; Random Forest 99.72%) benefits on the original deduplicated dataset outperform the other approaches. These models generalized well in spite of the low sample. Conversely, Naïve Bayes was outperformed by the former because its assumptions about feature independence was unstable under high dimensional environment.




Figure 5.3b Accuracy Comparison of PU Learning Models (Original Data)

Two-Step PU RF (99.72%) and Bagging PU (99.63%) are among the PU models that give robust performance with the original dataset. Elkanoto PU was slightly slower, probably because these required accurate estimation of class priors and conditionals probabilities, which become increasingly constraining with very few labeled samples.




Figure 5.3c Accuracy Comparison of Supervised Learning Models (Synthetic Data)

Because of the high accuracy the training of supervised models on the synthetic- augmented dataset resulted. There was a huge convergence of performance gap between models with ensemble methods within 98.6% and 98.87% score. This shows the progress of Gaussian-based synthetic augmentation in making the model robust and reducing variance.




Figure 5.3d Accuracy Comparison of Supervised Learning Models (Synthetic Data)

PU models also derived some metrics from synthetic augmentation. Bagging PU (98.73%) and Two-Step PU RF (98.69%) maintained their stability and competitiveness as their gap with supervised models was gradually reducing. This implies that synthetic data does not only promote supervised learning but also supports PU Learning strategy by increasing the positive-unlabeled training space.

This comparative evaluation strengthens synergy between traditional supervision and semi-supervised PU learning, especially with synthetic augmentation. The result presents a convincing argument for hybrid approaches in real-world IoV intrusion detection where the data from labeling is limited yet adaptability is important.

5.4 Prediction Comparison of best PU Model and best Supervised Model on Test Data

To study and compare the performance of the best models in supervised and PU learning, a thorough comparison of Two-Step PU Learning (Random Forest) and the Voting Classifier with respect to a number of predictions and classification performance metrics in their original and in synthetic datasets was performed.




Figure 5.4a Predicted Positives Two-Step PU vs Voting Classifier (Original Data)




Figure 5.4b Predicted Positives Two-Step PU vs Voting Classifier (Original Data)

With the original deduplicated dataset (3,588 records), both the Voting Classifier and Two-Step PU Learning were able to achieve high results in classification. From Figure 5.3.4a we see that the Voting classifier and Two-Step PU RF achieved F1-scores near to 0.92, with AUC scores of 0.972 and 0.974. respectively.

Even if equally accurate, the number of predicted positive findings (attacks) differed vastly. For example, Naive Bayes had 72 positives, but with a low F1-part of about 0.22, meaning high false positives, and, therefore, model overfitting. Bagging PU was the most conservative form, detecting only 8 positives, while Two-Step PU RF performs the best form because it equalizes precision and recall with 12 true positives predicted.

One of the great advantages of the model Two Step PU is in its use of Reliable Negatives (RNs). < These are benign samples in unlabeled training set that have a very low probability of being an attack. With a threshold of 0.2, the model detected 2,481 dependable negatives from 2,482 samples and this informatively increased the training quality of the model, despite having only twenty-nine labelled attack samples.

Further, Figure 6A bounds the overall F1-score vs. positives predicted for all models trained on the original dataset. This visualization shows that the best tradeoff between accuracy and correct attack detection was achieved by Two-Step PU RF and Voting Classifier while models such as Naive Bayes demonstrated poor generalization.




Figure 5.4c Predicted Positives Two-Step PU vs Voting Classifier (Synthetic Data)




Figure 5.4d Predicted Positives Two-Step PU vs Voting Classifier (Synthetic Data)

We then assessed the performance of the model in the synthetic-augmented dataset, which had a total of 77,380 records and was produced using Gaussian based sampling. The increased diversity and balance of the dataset greatly boosted all model performances.

As can be seen from Figure 5.3.4b, both Two-Step PU RF and Voting Classifier predicted a larger number of attack samples where the F1-scores were above 0.98. Two-step PU model picked out 27939 reliable negatives (same threshold of 0.2) out of 28310 unlabeled benign samples. This helped the model generalize better with more confidence whilst reducing false positives.

Figures 6A and 6B provide a comparative overview of F1-scores (bars) and number of predicted positives (lines) of all supervised and PU learning models for the original and synthetic datasets, respectively.

In Figure 6A, the trained models on the original dataset show high variance in the prediction, with the Voting Classifier and Two-Step PU RF producing stable results.

In the models trained on the synthetic data presented in figure 6B, there are higher and more consistent F1-scores and predictions, indicating the success of the augmentation.

These findings confirm the power of Two-Step PU Learning when used hand in hand with reliable Negative Selection and Synthetic Augmentation for intrusion detection in vehicular networks. Its durability in the face of scarce labels makes it adequate for practical use with scarce or missing annotations.

CHAPTER VI

CONCLUSION AND FUTURE WORK

6.1 Conclusion

The objective of this research journey is to contribute to improving cybersecurity in the Internet of Vehicles (IoV) through intelligent intrusion detection methods based on supervised and Positive-Unlabeled (PU) learning methods. Analyzing how these attacks affect the vehicular network manually is both a laborious and incomplete task, which we addressed through rigorous preprocessing, synthetically augmented data, and model evaluation to develop a comprehensive method that would be able to accurately discover DoS and spoofing attacks on vehicular CAN bus networks with high accuracy and adaptability.

Training these models on a balanced dataset gave them incredibly reliable results when it comes to using traditional supervised models like Random Forest, XGBoost, and Voting Classifiers. Furthermore, PU learning approaches like ElkanotoPU (WPb), BaggingPU (WPb++), PU-SVM (WpSVM), Two-Step PU (Wp+), etc. were proved in their relative applications in real world where labelled data is usually incomplete or unreliable. By including synthetic data, model performance saw an order of magnitude increase, as well as a massive reduction in false negatives (i.e., attack types in the minority of data) for those attack types in the minority of the data.

From a broader perspective, this project showcases the growing importance of semi-supervised strategies and synthetic data generation in the domain of vehicular cybersecurity. The findings further suggest that purely based on fully labeled sets may not be feasible in many of the IoV deployments, and need is the adaptive learning mechanism.

6.1 Key Takeaways

- When only positive attack samples are present, the problem of PU learning provides a scalable solution.
- Data augmentation on synthetic data balances the dataset and unlocks the power of ensemble and boosting methods.
- Regardless, XGBoost and Random Forest always had the best results in both supervised and PU pipelines.
- The first step is to take care of data cleaning and preprocessing—without removing duplicates, no, not normalizing features, the models could very well overfit or generalize poorly.

6.3 Limitations

Yet, the models achieved an extremely high performance, however there are still open spaces for further improvement.

- Gaussian assumptions for synthetic data generation do not always result in dependence on Gaussian feature distributions.
- Nevertheless, uncertainty can still exist in PU learning since assumptions on unlabeled data are made.
- These models were not tested for real time and scalability deployment and scalability in live vehicular systems.

6.4 Future Work

Future directions of this project are many exciting as they look ahead.

- Onboard Real-Time Deployment:

A logical continuation of this work is the direct employment of the trained models into onboard units (OBUs) or Electronic Control Units (ECUs). Real-time would require the implementation review of detection latency, memory overhead and model throughput in embedded environments; hence, would be a requirement to ensure IDS can respond quickly in live automotive operation.

- Execution Time Benchmarking with Timestamps:

Although concerns like accuracy and F1-score were at the heart of the script in focus, there were no explicit excerpts in which the execution time (timestamp) for training and inference specific times were discussed. Responses to future studies should note the number of training and testing periods of the models for per-model periods, which can be used to select models based not only on their predictive power but also on how quickly they compute. This is especially important in real-time automotive systems where the computation resources are limited.

- Deep learning models for both temporal and Sequence-aware tasks.

There is an inherent temporality associated with CAN traffic. Deep learning models like CNNs, RNNs, LSTMs, or Transformer architectures can be explored by future research to detect temporal patterns of attacks. Improvement of detection of time-sensitive anomalies or replay attacks that cannot be detected with static analysis using these models.

- Variations in the PU training configuration:

In the present study, PU Learning was executed by training and testing the algorithm based on a single 80-20 split of the training data and a predefined of 0.2 as a reliable Predefined threshold. But the performance of PU models depends on:

- Training/Testing ratios (e.g., 60-40, 70-30, 90-10)
- Class prior assumptions
- Probability thresholds for RN selection

These variations should be examined by future experiments as to how they affect recall, false positive rate, and reliability of negatives, particularly for the case of extreme class imbalance. As an example, RN stability can be evaluated across thresholds (e.g., 0.1 to 0.5) to find the optimal trade-off between precision and coverage. The tracking of the standard deviation of predicted positives across splits can be used as a metric for PU model robustness.

- Alternative Semi-Supervised Strategies:

In addition to PU learning, contrastive PU learning, Mix Match, and self-training using pseudo labeling on IoV environments should also be attempted. These methods can prove useful in cases when labels are scarce, or when the adversarial attack becomes out of the scope of current training data.

- Advanced Synthetic Generation with GANs:

Gaussian based augmentation did aid in imbalance, but future work can experiment with Generative Adversarial Networks (GANs) to produce high fidelity synthetic CAN messages. By being able to learn closer to the distribution of legitimate and malicious traffic, GANs may improve the diversity and the realism of training data.

- Cross-Domain and Multi-Vehicle Evaluation:

To provide a more real- world applicable result, it is important to test model generalization across diverse data from different vehicle types, manufacturers, and driving conditions. With more data coming in, this validation would become more robust across ECUs, CAN protocols and vehicular topologies.

- Lightweight IDS Models and Feature Pruning.

Feature selection techniques can be considered for selecting the minimum subset of predictors that can provide accurate classification. When combined with lightweight models such as MobileNet and Quantized XGBoost, this can then be applied on low power embedded platforms without performance loss.

Finally, this work provides a method for developing robust, accurate, and interpretable intrusion detection systems for connected vehicles. We show that it is possible to significantly augment the IoV security framework, through a hybrid approach that incorporates the fruits of classic supervision along with modern semi-supervised techniques as well as smart data augmentation, while opening the door to smarter, safer transportation ecosystems.

REFERENCES

- [1] Sharma, A., Kumar, R., & Singh, P. (2024). COMPARATIVE STUDY OF MACHINE LEARNING ALGORITHMS USING THE CICIOV2024 DATASET. *International Journal of Cybersecurity Intelligence & Cybercrime*, 7(1), 75–88.
- [2] Singh, A., & Verma, R. (2024). Evaluating Machine Learning Algorithms for Intrusion Detection: A Step Toward Securing Real-Time Big Data. *International Journal of Advanced Research in Computer Science*, 15(3), 150–165.
- [3] University of New Brunswick. (2024). CICIoV2024 Dataset. Canadian Institute for Cybersecurity. <https://www.unb.ca/cic/datasets/iov2024.html>
- [4] Liu, H., Wang, Y., & Zhang, T. (2023). Machine Learning Models for Intrusion Detection in CAN Networks: A Survey. *Journal of Automotive Cybersecurity*, 9(1), 45–66.
- [5] Elkan, C., & Noto, K. (2008). Learning classifiers from only positive and unlabeled data. *Proceedings of the 14th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 213–220.
- [6] Breiman, L. (2001). Random Forests. *Machine Learning*, 45(1), 5–32.
- [7] Chen, T., & Guestrin, C. (2016). XGBoost: A scalable tree boosting system. *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 785–794.
- [8] Cortes, C., & Vapnik, V. (1995). Support-Vector Networks. *Machine Learning*, 20, 273–297.
- [9] Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., ... & Duchesnay, E. (2011). Scikit-learn: Machine Learning in Python. *Journal of Machine Learning Research*, 12, 2825–2830.
- [10] Garcia, S., Luengo, J., & Herrera, F. (2015). *Data Preprocessing in Data Mining*. Springer.
- [11] He, H., & Garcia, E. A. (2009). Learning from Imbalanced Data. *IEEE Transactions on Knowledge and Data Engineering*, 21(9), 1263–1284.
- [12] Chawla, N. V., Bowyer, K. W., Hall, L. O., & Kegelmeyer, W. P. (2002). SMOTE: Synthetic Minority Over-sampling Technique. *Journal of Artificial Intelligence Research*, 16, 321–357.
- [13] Yang, X., Wang, J., & Zhang, Y. (2022). Intrusion Detection in the Internet of Vehicles: A Deep Learning-Based Approach Using CAN Bus Data. *IEEE Internet of Things Journal*, 9(7), 5403–5414.
- [14] Zhang, Y., & Zhao, M. (2023). Data Cleaning and Feature Engineering for Intrusion Detection Systems. *Proceedings of the IEEE Symposium on Cybersecurity and Analytics*, 112–118.

- [15] Rahman, M. M., & Islam, S. (2021). A Survey on Semi-Supervised Learning Techniques in Intrusion Detection. *ACM Computing Surveys*, 54(3), 1–35.
- [16] Saito, T., & Rehmsmeier, M. (2015). The Precision-Recall Plot Is More Informative than the ROC Plot When Evaluating Binary Classifiers on Imbalanced Datasets. *PLOS ONE*, 10(3), e0118432.
- [17] Kingma, D. P., & Ba, J. (2014). Adam: A Method for Stochastic Optimization. *arXiv preprint arXiv:1412.6980*.
- [18] Zhao, Z., & Salcic, Z. (2022). Evaluating PU Learning for Real-Time CAN Bus Intrusion Detection. *Proceedings of the International Conference on Cybersecurity and Resilient Systems*, 130–140.
- [19] Scikit-learn Documentation. (2023). <https://scikit-learn.org/stable/>
- [20] Matplotlib Developers. (2023). Matplotlib: Visualization with Python. <https://matplotlib.org>
- [21] Seaborn Developers. (2023). Seaborn: Statistical Data Visualization. <https://seaborn.pydata.org>
- [22] Python Software Foundation. (2023). Python Language Reference. <https://www.python.org>
- [23] Yuan, Z., Lu, Y., & Wang, Y. (2020). CAN Bus Intrusion Detection Using Ensemble Learning and Feature Selection. *IEEE Access*, 8, 152758–152769.
- [24] Alshamrani, A., Myneni, S., Chowdhary, A., & Huang, D. (2019). A Survey on Advanced Persistent Threats: Techniques, Solutions, Challenges, and Research Opportunities. *IEEE Communications Surveys & Tutorials*, 21(2), 1851–1877.
- [25] Jadhav, R., Raut, S., & Deshmukh, A. (2021). Supervised Machine Learning for Intrusion Detection: A Review. *International Journal of Security and Its Applications*, 15(2), 1–14.
- [26] Pan, S. J., & Yang, Q. (2010). A Survey on Transfer Learning. *IEEE Transactions on Knowledge and Data Engineering*, 22(10), 1345–1359.
- [27] Salama, M. A., Eid, H. F., & El-Sayed, A. (2020). A Hybrid Intrusion Detection Framework for Vehicular Ad Hoc Networks. *Journal of Network and Computer Applications*, 103, 119–132.
- [28] Wang, W., Sheng, Y., Wang, J., Zhu, M., & Zhang, Y. (2017). HAST-IDS: Learning Hierarchical Spatial-Temporal Features Using Deep Neural Networks to Improve Intrusion Detection. *IEEE Access*, 6, 1792–1806.
- [29] Cao, H., Wang, X., Chen, J., & He, L. (2020). Data Augmentation for Intrusion Detection Using GANs. *International Conference on Cloud Computing and Security*, 439–452.

- [30] Kim, H., Lee, H., & Kim, H. (2021). PU Learning-Based Intrusion Detection for IoV Environments. *IEEE Sensors Journal*, 21(3), 2344–2355.
- [31] Zhang, X., Jin, L., & Wang, Y. (2022). A Comparative Study of ML Algorithms for DoS Attack Detection in CAN Networks. *Computers & Security*, 114, 102588.
- [32] Roy, A., Niyogi, R., & Shukla, M. (2020). Data-Centric Approaches for IoT Intrusion Detection: Review and Recommendations. *Journal of Information Security and Applications*, 55, 102582.
- [33] Yu, R., & Liu, Z. (2021). Evaluating Semi-Supervised Methods for Cyberattack Detection in Vehicular Networks. *Future Generation Computer Systems*, 118, 36–47.
- [34] Subramanian, S., & Mishra, B. K. (2023). Feature Selection Techniques in IoT Intrusion Detection: A Comparative Analysis. *Computers, Materials & Continua*, 74(2), 2651–2667.
- [35] Ahmed, M., Mahmood, A. N., & Hu, J. (2016). A Survey of Network Anomaly Detection Techniques. *Journal of Network and Computer Applications*, 60, 19–31.
- [36] Shone, N., Ngoc, T. N., Phai, V. D., & Shi, Q. (2018). A Deep Learning Approach to Network Intrusion Detection. *IEEE Transactions on Emerging Topics in Computational Intelligence*, 2(1), 41–50.
- [37] Canadian Institute for Cybersecurity. “CICIoV2024 Dataset.” University of New Brunswick, 2024. [Online]. Available: <https://www.unb.ca/cic/datasets/iov2024.html>